

# How to calculate an election result in your bedroom

by Michael Sproul



# Australia!



- Population: 24 million
- 8 states and territories
- Senate composition is decided *per state*



# STV Preferential Voting

- Voters number candidates in order of preference
- Algorithm elects popular candidates, knocks out unpopular ones
- Lots of candidates, lots of ballots and complicated rules!

# The AEC's dirty secret

The Australian Electoral Commission keeps their code secret!

... but all the data is available as CSV!

... and the voting algorithm is defined in legislation!

# Why Rust?

- Memory usage  
 $10^2$  candidates  $\times$   $10^7$  voters  $\times$  4 bytes  $\approx$  1GB RAM
- Speed  
As fast as C/C++ without segfaults
- Types  
Easy refactoring, controlled mutability, ADTs (enums)



# Naive Number Crunching

Example:

Candidate A has 100,000 votes

Quota = 80,000 votes

Elect Candidate A and transfer each vote to the next round with value  $20,000/100,000 = 1/5$

Logically, 4/5ths of each vote have been “used up”

# Naive Number Crunching

**What do we do with votes that are already fractional?**

Naive Solution: Multiply their value by the new transfer value

E.g.

Candidate B has 80,000  $\frac{1}{5}$  votes

Apply transfer value of  $(1/5)/80,000 = 1/400,000$

Votes from previous round now have value  $1/5 * 1/400,000 = 1/2,000,000$



# Naive Number Crunching

Uh oh!

- Irreducible fractional votes with **3000 digits in numerator and denominator**
- Standard `num` crate too slow, switched to `gmp` and `ramp`
- 20+ hours to run, 8+ gigabytes of memory used!



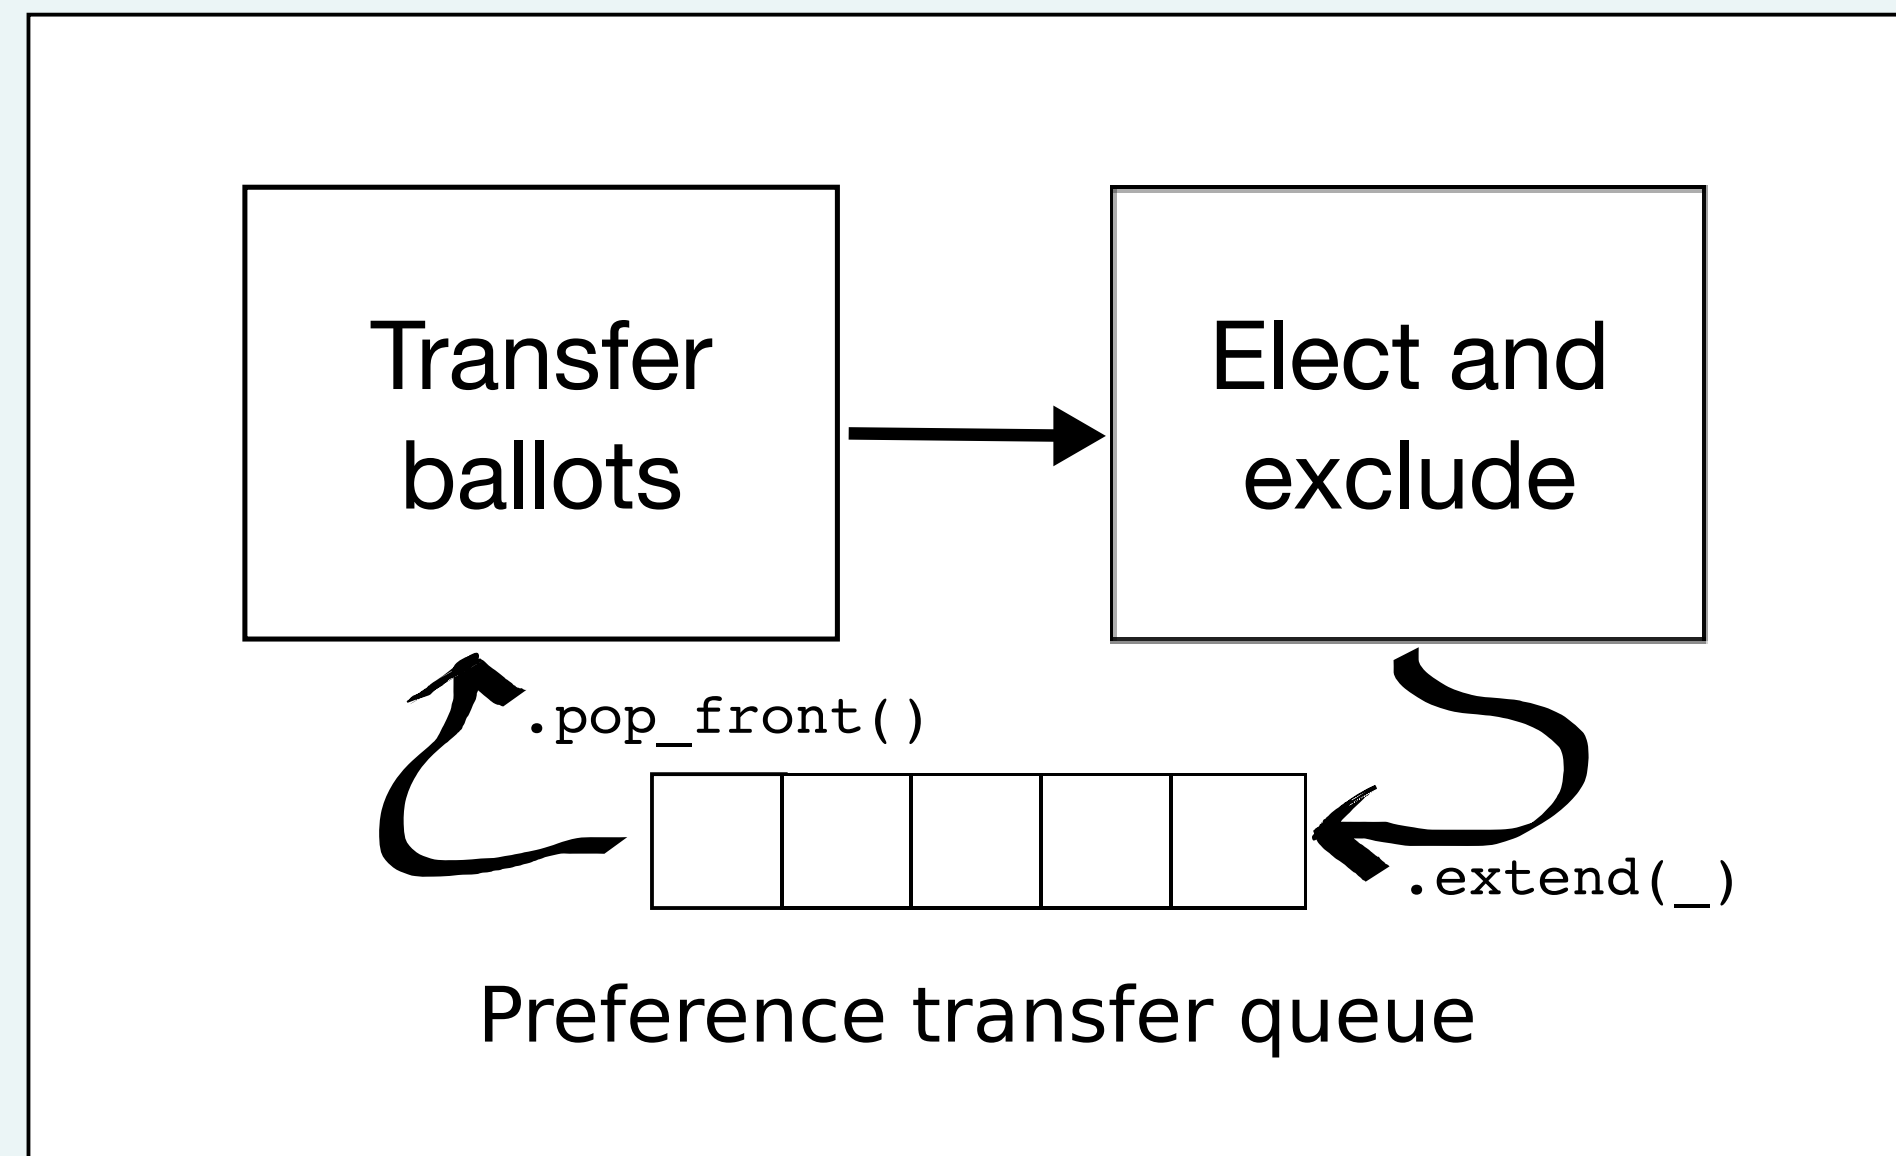
# Reading the Legislation

- Round vote counts to integers
- “Aggregate” transfer values to avoid multiplying them:  
$$t = (\text{num\_votes} - \text{quota}) / (\text{num\_ballots})$$
- Downside: more complex, more sensitive to order of operations

# System Architecture



Main algorithm, round  $i$



# Core Data Structures

```
pub struct VoteMap<'a> {  
    info: HashMap<CandidateId, VoteInfo<'a>>,  
    candidates: &'a CandidateMap  
}  
  
struct VoteInfo<'a> {  
    votes: Int,  
    ballots: TransferMap<'a>,  
    eliminated: bool,  
}  
  
pub type TransferMap<'a> = BTreeMap<Frac, Vec<&'a mut Ballot>>;
```

# Results

- Replicated the official result for all 8 states and territories
- For the largest state (NSW):
  - 4,705,270 votes
  - 20 seconds to run
  - ~1GB peak memory usage
- 1021 lines of code

# Links

- My code: [github.com/michaelsproul/aus\\_senate](https://github.com/michaelsproul/aus_senate)
- Another implementation: [github.com/grahame/dividebatur](https://github.com/grahame/dividebatur)

# Image Credits

- *Rust Logo* by Mozilla, CC-BY <https://www.rust-lang.org/en-US/legal.html>
- *Australia Map* by Lokal\_Profil, CC-BY-SA-2.5 [https://commons.wikimedia.org/wiki/File:Australia\\_map,\\_States.svg](https://commons.wikimedia.org/wiki/File:Australia_map,_States.svg)
- *The Queen* by UK Ministry of Defence, OGL v.3 [https://www.defenceimagery.mod.uk/fotoweb/archives/5000-Current%20News/Archive%20\(Navy\)/RoyalNavy/2015/March/45158590.jpg](https://www.defenceimagery.mod.uk/fotoweb/archives/5000-Current%20News/Archive%20(Navy)/RoyalNavy/2015/March/45158590.jpg)
- *CSV Icon* made by Freepik from [www.flaticon.com](http://www.flaticon.com)

Thank you!